

Machine Learning

Paria Khalili 810801065

Phase 1 Report

Introduction to Language Identification

Spoken Language Identification (SLID) is a computational process where a system categorizes an input acoustic signal into one or more predefined language categories. Unlike Speech-to-Text (ASR), which cares about *what* is being said, SLID cares about *how* the phonetic and rhythmic structures are organized.

Routing in Multilingual Systems: For global voice assistants (Alexa, Siri), SLID acts as a "gatekeeper." If the system can't identify the language in the first 500ms, it cannot load the correct acoustic model for transcription, leading to total failure.

Emergency Services: In international call centers, SLID is used to automatically route callers to a translator.

Low-Resource Scenarios: SLID helps in identifying dialects that might not have a written form but need to be categorized for digital archiving.

In Machine Learning, the goal is to create a model that learns "Invariance", the ability to recognize the language regardless of where or how it was recorded.

Acoustic Variance: The "Channel Effect"

When you record in a Controlled Environment, you are minimizing the "Channel Effect." The channel is everything between the speaker's mouth and the digital file (the room, the microphone, the preamp).

- **Controlled Data:** The frequency response is usually "flat." This means the energy levels across all frequencies (20Hz to 20kHz) are preserved accurately.
- **Audiobook Data:** Audiobooks often undergo Dynamic Range Compression and Equalization (EQ) during post-production to make the narrator's voice sound "warm" or "intimate."

For an LID model, this is a challenge. If the German audiobook has a heavy bass boost (common in professional narration), a naive model might mistake "bass-heavy audio" as a feature of the German language itself. This is why practitioners use Cepstral Mean and Variance Normalization (CMVN) to "center" the features and strip away the channel signature.

The "Cocktail Party Effect" vs. "Narrative Clarity"

While audiobooks are "real-world," they are a specific *kind* of real world.

- **Controlled Data:** Perfect isolation. No "overlapping speech."
- **Audiobooks:** Usually a single narrator, but they may perform character voices. A female narrator might lower her pitch to play a male character or use a high-pitched, breathy tone for a child.
- **True Real-World (Street recordings):** Background noise, multiple people talking, traffic.

Audiobooks are the perfect "Stress Test" for Speaker Independence. Because one narrator (Nadine Dela) is playing multiple roles or tones, the model is forced to learn that "German" exists in both the high-pitched character voice and the low-pitched narrative voice.

This prevents the model from simply memorizing the speaker's "pitch" as the identifying feature.

In a lab setting, speakers are often asked to speak at a "nominal" rate (about 150 words per minute). In audiobooks, the Speech Rate Variability is massive.

Models trained on controlled data often fail in the real world because they "overfit" to the silence or the microphone quality. Audiobooks—like the one you are using—are considered "Silver Standard" data. They provide the clarity of a studio but the linguistic complexity of natural storytelling.

Closed-Set vs. Open-Set LID:

Closed-Set LID

- Concept: The system assumes the input *must* belong to one of N known languages (e.g., your project might only look at German, English, and French).
- Math: Uses a standard Softmax function where the sum of all probabilities equals 1.0 (100%).
- Weakness: If you provide a clip in Japanese to a Closed-Set model trained only on German and English, it will "force" a choice, often giving a high-confidence wrong answer (a "False Positive").

Open-Set LID (The Real-World Reality)

- Concept: The system must recognize N known languages AND a "Reject" or "Unknown" category.
- Math: Often uses Thresholding or Out-of-Distribution (OOD) Detection. Instead of a simple Softmax, it may use a Sigmoid layer for each language or calculate the "Entropy" of the prediction. High entropy = "I don't know this language."

-
- Practicality: Much harder to train because the "Unknown" category is infinite (it represents every other language on Earth).

Implementation Approaches for Each Paradigm

Implementation has evolved through three distinct "Eras":

A. The Phonetic Approach (PRLM)

- Method: Phone Recognition followed by Language Modeling.
- Logic: The system converts audio into "Phonemes" (the smallest units of sound). It then calculates the probability of certain sound sequences. For example, the sound "pf" is very common in German but rare in English. If the system sees many "pf" phonemes, it flags the language as German.

B. The Statistical Approach (I-Vectors)

- Method: Universal Background Models (UBM) and Factor Analysis.
- Logic: This approach models the "Total Variability" in a voice. It compresses the audio into a fixed-length vector called an I-vector. This vector represents the speaker's characteristics and the language characteristics simultaneously.

C. The Deep Learning Approach (X-Vectors & Transformers)

- Method: Using CNNs (Convolutional Neural Networks) or Transformers.
- Logic: * X-Vectors: These are the current industry standard. A Deep Neural Network (DNN) is trained to recognize speakers, and the "embeddings" (internal layers) are extracted to identify the language.

-
- Spectrogram Analysis: The model treats the audio like an image, looking for visual patterns in the Mel-Spectrogram that represent the "texture" of the language.

In Machine Learning for speech, we don't just look at the frequencies (the "notes"); we look at the temporal structure (the "rhythm"). German is fundamentally different from languages like French or Spanish because of how it handles time.

Stress-Timed (German) vs. Syllable-Timed (French/Spanish)

To understand why this is a "deep" observation, you must distinguish between these two rhythmic classes:

- Syllable-Timed (Isochrony): In languages like Spanish, every syllable takes roughly the same amount of time to say. It sounds like a "machine gun" (Ta-ta-ta-ta).
- Stress-Timed (German/English): The time between stressed syllables is equal. To maintain this rhythm, German "squashes" the unstressed syllables in between. This leads to complex consonant clusters and varied vowel lengths.

The Impact on "Rhythm Metrics"

In SLID, practitioners use specific mathematical formulas to quantify this rhythm. Because you chose a German audiobook, your model will encounter these specific "Rhythm Metrics":

A. Delta measures (Delta C and Delta V)

This calculates the standard deviation of the duration of Consonant (C) and Vowel (V) intervals.

- In German: Because we squash unstressed vowels and have long, complex consonant clusters (like in the word "*Schlecht*"), the Delta C is very high. The consonants vary wildly in length compared to a language like Italian.

B. The Varco Measures (VarcoC and VarcoV)

This is the Delta measure normalized by the average speaking rate. It ensures that if the narrator speaks faster during an intense part of the book, the "rhythm score" stays consistent.

C. nPVI (Pairwise Variability Index)

This is the most famous metric in LID. It looks at the difference in length between two *adjacent* vowels.

- German Result: High nPVI. A stressed, long German vowel is often followed by a very short, reduced "schwa" sound (like the "e" in "*haben*"). This "long-short" pattern is a dead giveaway for German.

When you feed your audio clips into a model, it isn't just listening to the narrator's voice; it is learning these Duration Patterns:

1. Phonotactic Constraints: German allows "consonant heavy" segments. Your model will see high energy in the high-frequency bands (where "s", "sh", and "t" live) for longer durations than in a Romance language.
2. Vowel Reduction: The model learns that German vowels aren't always "pure." They change shape and shorten depending on their position, which is a key feature of stress-timing.

By selecting German audiobook data, we provide the system with a rich source of stress-timed isochrony. Unlike syllable-timed languages, German exhibits high variance in consonant interval durations (Delta C) and significant pairwise variability in vowel length (nPVI). These rhythmic signatures serve as robust non-spectral features that enhance the model's ability to discriminate German from languages with more uniform rhythmic structures.

Challenges in Language Identification

Phonetic Similarities Between Related Languages

This is the challenge of "Inter-class Similarity."

The Problem: Languages from the same family (e.g., the Germanic family including German, Dutch, and English) share a massive percentage of their phonemes (sounds).

- The ML Impact: When two classes are close together in the feature space, a standard classifier (like a Softmax CNN) struggles to draw a clean "decision boundary." For example, the "ach-laut" (/x/) in German is also present in Dutch.
- Deep Insight: For your report, mention that models must rely on Phonotactics—not just the *sounds* themselves, but the *rules* of how sounds follow one another. German allows certain consonant clusters (like "pf" or "tsch") that are illegal in other phonetically similar languages.

Accent and Speaker Variability

This is the challenge of "Intra-class Variability."

- Accent: Non-native speakers (L2 speakers) often carry the "phonetic blueprint" of their first language into their second. An LID system might misclassify a French person speaking German as "French" because it over-focuses on the accent rather than the language structure.
- Speaker Identity: Every person has a unique vocal tract. This introduces "Indexical Information" (age, gender, health).
- The ML Risk: If a model is not properly "normalized," it might learn to identify the narrator (Nadine Dela) rather than the language (German).

Short-Duration Speech Segments

The instructors have asked you for 1-minute clips, which is actually "luxurious" in the ML world.

- The 3-Second Rule: Most production systems (like Google Assistant) need to identify the language in under 3 seconds.
- The Problem: Feature extraction becomes "noisy" with less data. A 1-second clip might only contain one or two words, which may not be representative of the language's average rhythm or phonetic distribution.
- Accuracy Drop: Studies show that LID accuracy drops exponentially as the clip duration falls below 5 seconds because the Prosodic and Rhythmic features (like the stress-timing we discussed) require a longer time window to become statistically visible.

Review of Research Directions and Solution Strategies

1. Multi-Task Learning (MTL)

Instead of just training the model to predict "Language," researchers train it to predict Language + Gender + Speaker ID simultaneously.

- Why it works: By forcing the model to identify the speaker and gender in separate "auxiliary tasks," the main "Language Identification" branch is forced to focus only on the linguistic features that are *independent* of who is talking.

2. Self-Supervised Learning (SSL) Representations

The current state-of-the-art uses models like Wav2Vec 2.0 or Whisper.

- Strategy: These models are "pre-trained" on 50,000+ hours of unlabeled audio. They learn the universal "physics of speech" before they ever see a language label.
- Benefit: These models are much more robust to short segments because they already "understand" what a phoneme looks like, even in a 0.5-second clip.

3. Metric Learning (The "Triplet" Strategy)

As we discussed earlier, using Triplet Loss helps with phonetic similarity.

- Strategy: The model is trained to minimize the distance between "German Clip A" and "German Clip B" while maximizing the distance to "Dutch Clip C."
- Result: This creates a "Latent Space" where even very similar languages are pushed apart into distinct clusters.

4. Sequence Modeling (LSTMs and Transformers)

- Strategy: Using Long Short-Term Memory (LSTM) networks or Transformers to capture "Long-Range Dependencies."
- Logic: While a 1-second clip is hard, a Transformer can remember a sound it heard at the beginning of a sentence and compare it to a sound at the end. This allows the model to "calculate" the rhythm (stress-timing) across the entire 1-minute clip you provided.

In this project, the use of 60-second segments from audiobooks partially mitigates the 'short-duration' challenge. However, by selecting a single narrator performing multiple character voices, we are actively engaging with the 'Speaker Variability' challenge, forcing our feature extraction (like MFCCs) to be robust against changes in pitch and tone within the same language class.

Audio Data Characteristics and Selection

In Spoken Language Identification, we are trying to capture the Statistical Identity of a language. Unlike an image of a cat (which stays a cat regardless of the background), a speech signal is highly "volatile."

- **Feature Contamination:** If you select audio with heavy background music or white noise, the model might accidentally learn the "noise profile" of the recording equipment instead of the language. In your German audiobook, the "silence" between sentences is clean, ensuring the model's Signal-to-Noise Ratio (SNR) is high.
- **Acoustic Overlap:** If the selection includes multiple people talking at once (babble noise), the spectral features (like MFCCs) will be "smeared." This makes it impossible for the model to identify the specific formant transitions that characterize German phonemes.
- **Data Bias:** If you only select very fast speech, the model might fail on slow, dramatic narration. Proper selection ensures a balanced temporal distribution.

• Desirable Properties of Speech Samples for LID

To build a robust system, the audio you extracted from *Until the Last Heartbeat* should ideally possess the following properties:

-
1. **Phonetic Richness:** The segments should contain a wide variety of phonemes. Audiobooks are perfect for this because literary text uses a much broader vocabulary than daily conversation.
 2. **Stationarity (Short-term):** Speech is a "quasi-stationary" signal. We want segments where the voice is stable enough for 20ms-30ms "frames" to capture a clear spectral shape.
 3. **Linguistic Purity:** The sample should be free of "Non-Linguistic Vocalizations" (coughs, sneezes, or long periods of heavy breathing) that don't contribute to language patterns.
 4. **Isochronous Consistency:** As we discussed, German is stress-timed. A desirable sample is one that clearly demonstrates this rhythm—where the stressed "beats" of the German sentences are audible and measurable.

• The Role of Listening-Based Qualitative Evaluation

While "Technical Analysis" (looking at waveforms or SNR levels) is important, the human ear is still the most sophisticated "feature extractor" available.

Why we must listen (Qualitative):

- **Contextual Slicing:** A computer might cut a segment based on silence, but a human ensures the cut happens at a logical sentence boundary. As per your project requirements, cutting mid-phrase destroys the "Language Model" (the probability that word B follows word A).
- **Emotion vs. Language:** A human listener can detect if the narrator is using an "accented" character voice. If the narrator mimics a French accent while speaking German, the technical features will shift toward French. Qualitative evaluation allows us to discard these "confusing" samples to keep the training data "pure."
- **Detecting Artifacts:** Small digital clicks, "plosives" (the "popping" sound of 'p' or 'b'), or page-turning sounds might show up as high-energy spikes on a graph, but a human can instantly identify them as "non-speech" and decide if the clip is still usable.

Feature Extraction Techniques

Feature extraction is the process of converting a raw audio signal into a compact, informative representation that discards redundant information (like background noise or volume) while retaining phonetic and prosodic cues essential for identifying a language.

Mel-Frequency Cepstral Coefficients (MFCC)

- **How it Works:** MFCCs mimic human hearing. The signal is framed, windowed, and passed through a Fast Fourier Transform. The power spectrum is mapped onto the Mel Scale (logarithmic spacing of frequencies). Taking the Log and then the Discrete Cosine Transform (DCT) results in "cepstral" coefficients.
- **Suitability for LID:** Highly suitable. MFCCs capture the "envelope" of the short-term power spectrum, which represents the shape of the vocal tract—a key differentiator in phonetic production across languages.
- **Advantages & Limitations:** * *Pros:* Effective at capturing phonetic content; computationally efficient.
 - *Cons:* Sensitive to additive noise and channel distortion (e.g., different microphones).
 - *Reference:* Davis & Mermelstein (1980).

Fast Fourier Transform (FFT)

- **How it Works:** A mathematical algorithm that computes the Discrete Fourier Transform (DFT), converting a signal from the time domain to the frequency domain.
- **Suitability for LID:** Low. Raw FFT output is too high-dimensional and contains too much phase information and noise to be used directly for classification. It serves as a foundational step for other features.
- **Advantages & Limitations:** * *Pros:* Complete representation of the signal's frequency content.

-
- *Cons:* Linear frequency scaling does not align with human perception or linguistic importance.

Log-Mel Spectrogram

- How it Works: Similar to MFCCs, but it skips the final DCT step. It applies a Mel-filterbank to the power spectrum and takes the logarithm.
- Suitability for LID: Excellent, especially for Deep Learning (CNNs). It preserves spatial relationships between frequencies that MFCCs (which decorrelate data) discard.
- Advantages & Limitations: * *Pros:* Preferred for visual-based neural networks (Spectrograms as images).
 - *Cons:* Higher dimensionality than MFCCs.

Spectral Centroid

- How it Works: Indicates where the "center of mass" of the spectrum is located. It is calculated as the weighted mean of the frequencies present in the signal.
- Suitability for LID: Moderate. It can distinguish between "bright" sounding languages (higher frequencies) and "darker" sounding ones, but cannot identify specific phonemes.
- Advantages & Limitations: * *Pros:* Simple measure of spectral sharpness/brightness.
 - *Cons:* Very limited information; usually only used as a supplementary feature.

Chroma Features

- How it Works: Projects the entire spectrum into 12 bins representing the 12 distinct semitones (octaves) of the musical octave.
- Suitability for LID: Low to Moderate. Primarily used in music analysis, but can capture "tonality" in tonal languages (e.g., Mandarin vs. English).
- Advantages & Limitations: * *Pros:* Robust to changes in timbre.
 - *Cons:* Often irrelevant for non-tonal language identification.

Spectral Contrast

- How it Works: Measures the energy difference between the peaks and valleys in the frequency spectrum across different sub-bands.
- Suitability for LID: Moderate. It helps distinguish between harmonic (vowels) and non-harmonic (fricatives) components.
- Advantages & Limitations: * *Pros*: Better performance in noisy environments than MFCCs alone.
 - *Cons*: Computationally more intensive.

Zero-Crossing Rate (ZCR)

- How it Works: The rate at which the signal changes from positive to negative or vice versa (crosses the zero axis).
- Suitability for LID: Supplementary. High ZCR often indicates unvoiced speech or percussive sounds.
- Advantages & Limitations: * *Pros*: Extremely easy to calculate.
 - *Cons*: Highly sensitive to low-frequency noise.

Linear Predictive Coding (LPC)

- How it Works: Predicts future samples of a signal based on a linear combination of past samples. It estimates the parameters of a digital filter that models the vocal tract.
- Suitability for LID: Good for low-bitrate environments. It characterizes the resonance (formants) of the speaker.
- Advantages & Limitations: * *Pros*: Good for modeling the source-filter theory of speech.
 - *Cons*: Assumes a linear model, which speech isn't perfectly; less robust than MFCCs.
 - *Reference*: Makhoul (1975).

Perceptual Linear Prediction (PLP)

- How it Works: Combines LPC with psychophysical concepts (equal loudness pre-emphasis and intensity-to-loudness compression).
- Suitability for LID: High. It is designed to eliminate speaker-specific information while retaining linguistically relevant information.
- Advantages & Limitations: * *Pros*: Often outperforms MFCCs in speaker-independent tasks.
 - *Cons*: More complex to implement.
 - *Reference*: Hermansky (1990).

Similarity Learning

Definition

Similarity Learning is an area of supervised machine learning where the goal is to learn a distance function (metric) that measures how similar or related two objects are. In LID, instead of classifying a clip directly as "French," the model learns to project audio into an embedding space where clips of the same language are close together and clips of different languages are far apart.

Common Loss Functions

- Contrastive Loss: Takes a pair of inputs. If they are from the same language, the loss minimizes the distance between them. If they are from different languages, it maximizes the distance up to a certain "margin."
- Triplet Loss: Takes three inputs: an Anchor (e.g., German), a Positive (another German clip), and a Negative (e.g., Spanish). The loss ensures the Anchor is closer to the Positive than to the Negative by a specific margin.

Application to Language Discrimination

Similarity learning is particularly powerful for Open-Set LID. In traditional classification, a model can only identify languages it was trained on. With similarity learning, a model can be shown a new, unknown language and determine it is "different" from all known languages because its embedding falls into a new cluster. It allows for "Language Fingerprinting" where a short sample can be compared against a database of known language clusters to find the best match.